

Omulimu: Telegram coaching bot in n8n

Test project proposal · \$60 fixed, 48 hours

1. What to try on the live bot

Messages to try "sold 120k, spent 40k" · "how do I price my chapati?" · "I feel so stressed" · /help · a sticker

2. How a message flows

Telegram Trigger → Save to **messages** (ON CONFLICT DO NOTHING) → LLM classifier (last 15 messages)
→ Router (open session + wellbeing-first rule) → Numbers | Coach | Care | Help brain
→ **Send Reply** sub-workflow (claims the outbound row, 60-word cap, retry on 429) → Telegram
LLM fails 3 times → error branch: fallback reply, **error_log** row, message marked failed
Anything unhandled → Error Trigger workflow → alert in the admin chat

3. Technology

Layer	Choice	Why
Orchestration	n8n 1.123 (self-hosted, Docker)	7 workflows: a router, 4 brains, Send Reply and the error alert. They are versioned JSON with fixed IDs, so imports overwrite in place.
Intent classifier	gpt-4o-mini via OpenRouter, temperature 0, JSON mode	Returns one of 4 intents plus the amounts and language. The output is checked against a schema and repaired once if invalid. An OpenAI Decisions API mode is optional.
Brains	Numbers: plain code Coach: LLM, max 60 words Care: fixed templates Help: menu	Only the Coach writes free text. Numbers and Care use no LLM, so their sums and safety wording can't drift.
Database	Postgres (Supabase), 6 tables	Unique Telegram update id = no duplicates. The schema has users, messages, sessions, routing_history, daily_records and error_log, and is safe to re-run.
Hosting	Render web service (always on)	Keeps the webhook live and survives restarts: all n8n state lives in Postgres, and secrets live in environment variables, not in the exports.
Quality	24 unit tests, 55-check end-to-end suite	Covers duplicates, parallel updates, LLM failure, a 429 from Telegram, the word cap, buttons and the Care session.

4. Costing

Project fee

Item	Price
Build, test and hand over: n8n export (7 workflows), schema.sql, prompts and templates, setup README, test walkthrough	\$60 fixed
Delivery	Within 48 hours of the contract starting
Fixes for anything that fails the acceptance tests in your brief	Included

Running cost after handover (what you pay the providers each month)

Service	Plan	Monthly
Render (n8n host)	Starter. The free plan sleeps and misses Telegram messages.	\$7
Supabase (Postgres)	Free tier	\$0
Telegram Bot API	No charge	\$0
LLM (gpt-4o-mini)	Pay per use, about \$0.30 to \$0.60 per 1,000 messages	under \$1 at pilot volume
Total		about \$8 / month

The LLM figure is an estimate: about 1,500 input and 100 output tokens per classification, plus one Coach call for business questions. OpenRouter's dashboard shows the real spend per day. Already on a VPS? The same Docker setup runs there, and Render isn't needed.

5. What you receive

- The n8n workflow export (JSON, no secrets inside) and a credentials template
- **schema.sql** for an empty Postgres, and **prompts.md** with the classifier prompt, Coach prompt and every reply template
- A README with setup steps, how to add a fifth intent, known limits, and open questions for you
- Test results for your acceptance tests, plus SQL checks you can run yourself